

Summary of Last Class

- Pros & cons of monolith
- Break a Monolith  → SOA
- Pros & cons of microservices
stated with solutions

Cons

- ① Debugging across services is hard - Observability
- ② Consistency " " " " - Distributed Transactions
 - ↳ 2PC
 - ↳ Saga
 - ↳ Orch -- Chora.
- ③ Communication is hard - Event driven architecture - CQRS
- ④ Robustness - Circuit breaker

Observability

Solution / code is never perfect - always be some bugs / corner cases)
errors (n/w / disk) / requirements change

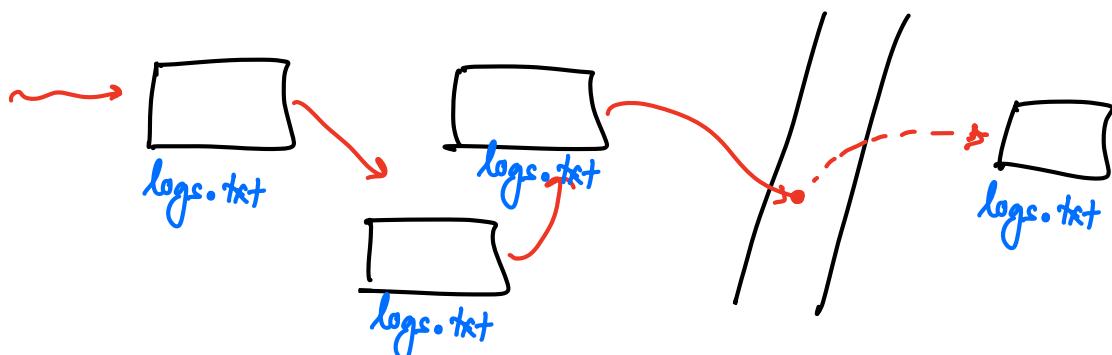
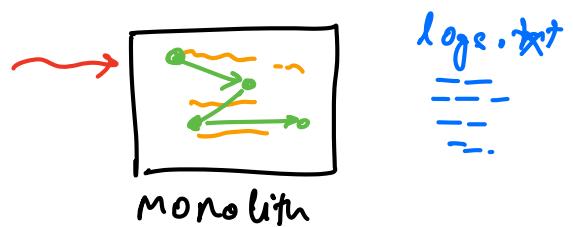
Errors are okay as long as they don't pass silently.

- ① Detection How will I get to know that something is wrong?
- ② Diagnosis Identify where the problem occurred
machine / service
& identify what caused it

③ Verification Once fixed, how do I verify that the fix works?

Diagnosis

Tracing a request



① When a request arrives at backend (LB) you can generate a unique correlation-id / request-id / trace-id / msg-id ...

Every time you take an action based on this request or forward this req

You MUST pass this correlation id

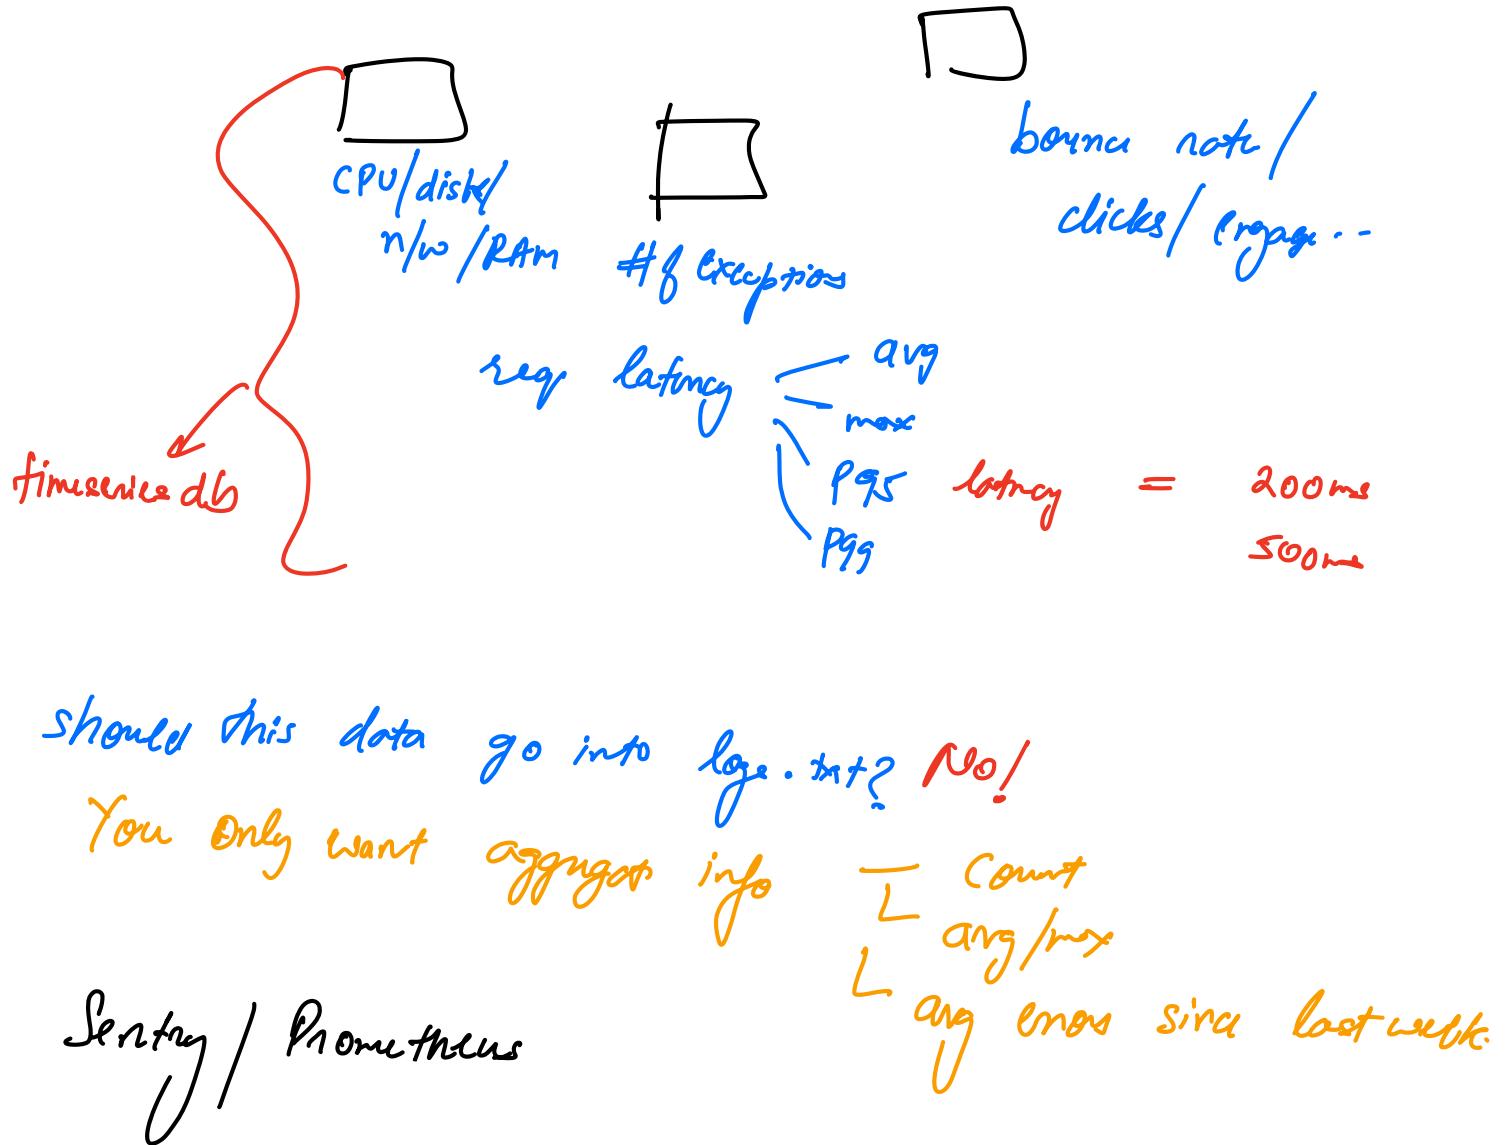
②

Aggregating Logs & searching



Defection & Verification

- trace certain metrics

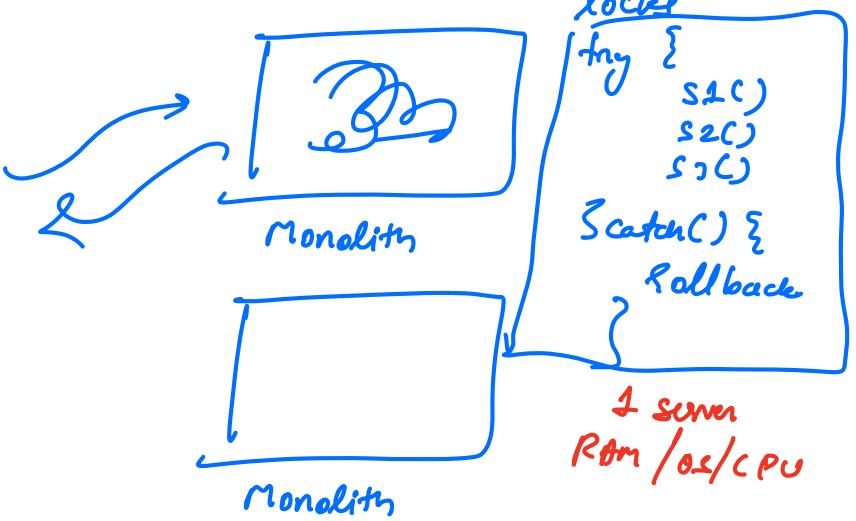
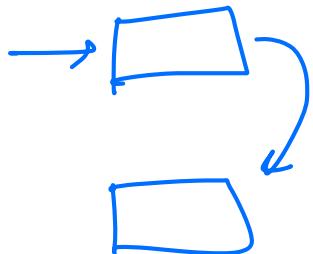


Alerts
detection

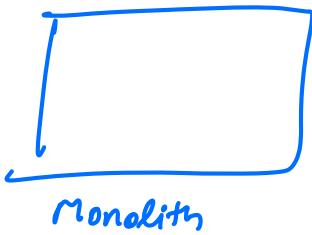
if # of exception / hour > 1000
└ send sms/email → ---

Verify
have # of exception metric improved?

Consistency



Microservices - any transaction might span multiple servers - v. hard
~~CAP~~
PACELC



Challenge: Consider two services A & B

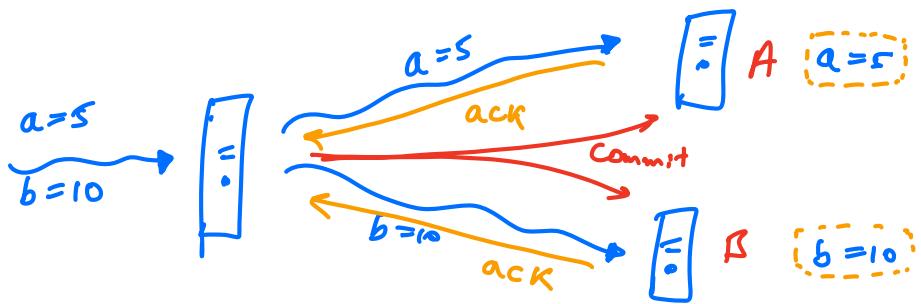
you wish to make a transaction that has multiple steps
 this transaction requires steps to be executed in both
 A & B

has to be atomic - either full failed / fully successful

How to do this reliably?

①

Two Phase Commit



① Prepare

Take note of this info
 do NOT commit it
 Plan acknowledgement
 Once ack you're NOT
 allowed to go back

② Commit

you had promised
 (Ack) some data earlier
 Please commit it

When these servers (A & B) save data & give ack they
MUST acquire locks. — No other reads/writes are allowed.
on Row/Key

v. Poor performance

∴ the state diagram is v. complex & has high depth
v. poor latency.

∴ 2PC is considered an anti-pattern in microservice

Pattern : Common & robust solution that is widely accepted as correct

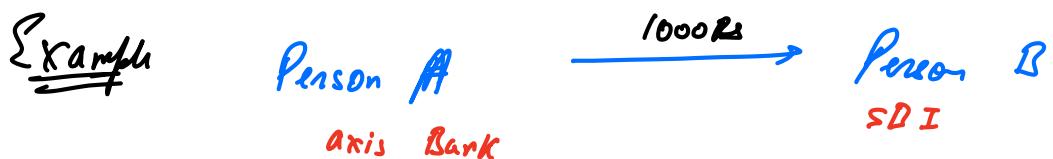
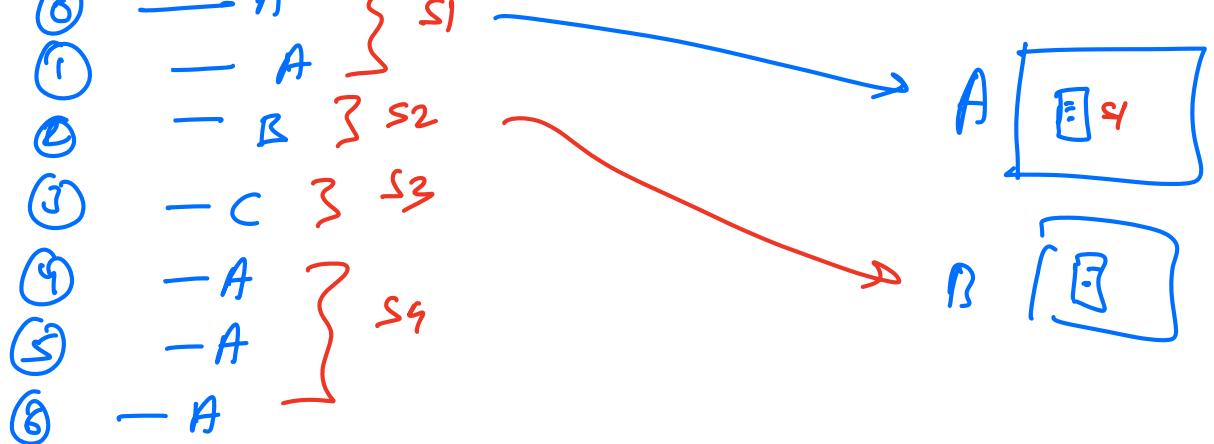
anti-pattern : common but flawed solution

② Saga Patterns

Break down a large & complex transaction into smaller 'Sagas'

→ involves lots of
steps that span across
multiple services

→ each saga should be 'local' to a specific microservice
You should not need to talk to multiple microservices
in any saga



- ① ensure A has sufficient balance } s_1 Axis
- ② deduct 1000 Rs from A }
- ③ add 1000Rs to B } s_2 SDI

- execute s_1
 - └ fail — show error
 - └ succeed
 - └ execute s_2
 - └ succeed — success(done)
 - └ fail — rollback ALL earlier steps

How to ensure that sagas are executed reliably?

- └ Orchestration — Music
- └ Chronography — Dance

Orchestration

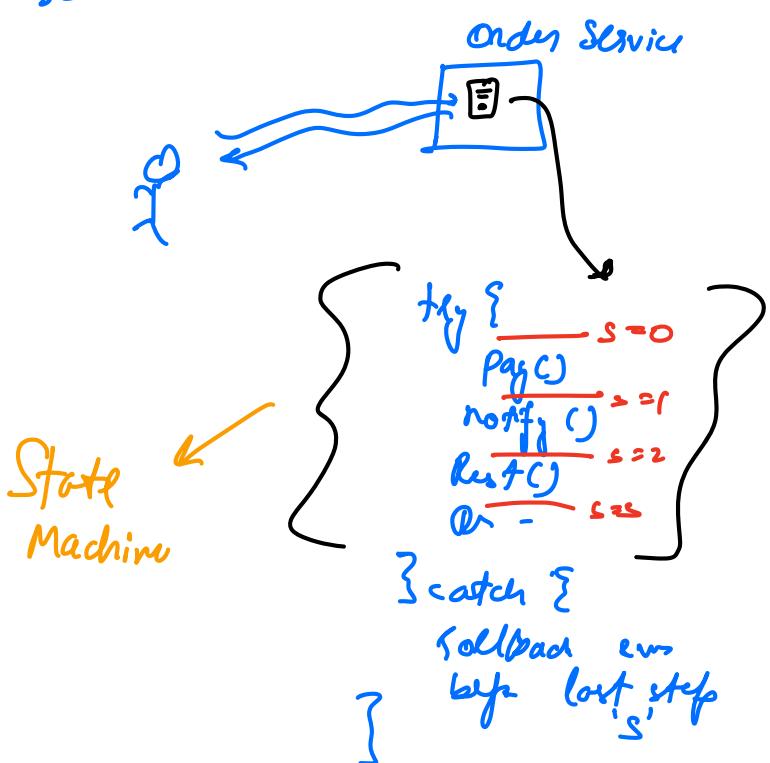
- central authority that oversees the execution of Sagas

a dedicated microservice to handle the transaction.



Swiggy - order food

- ① Order placement — SQL
- ② Payment
- ③ Notification
- ④ Restaurant
- ⑤ Delivery Assignment
- ⑥ Notification



Pros

Simple to understand

Cons

- Not simple to code — state machine
- Spof — even if we have multiple servers for orchestrator if one orchestrator server dies it becomes v.v. difficult to reproduce the faulty state & correct it
- detect? State is stored in DB

Choreography

- every service knows what it does
- each service is ONLY concerned about itself!!



Inherently Event Driven

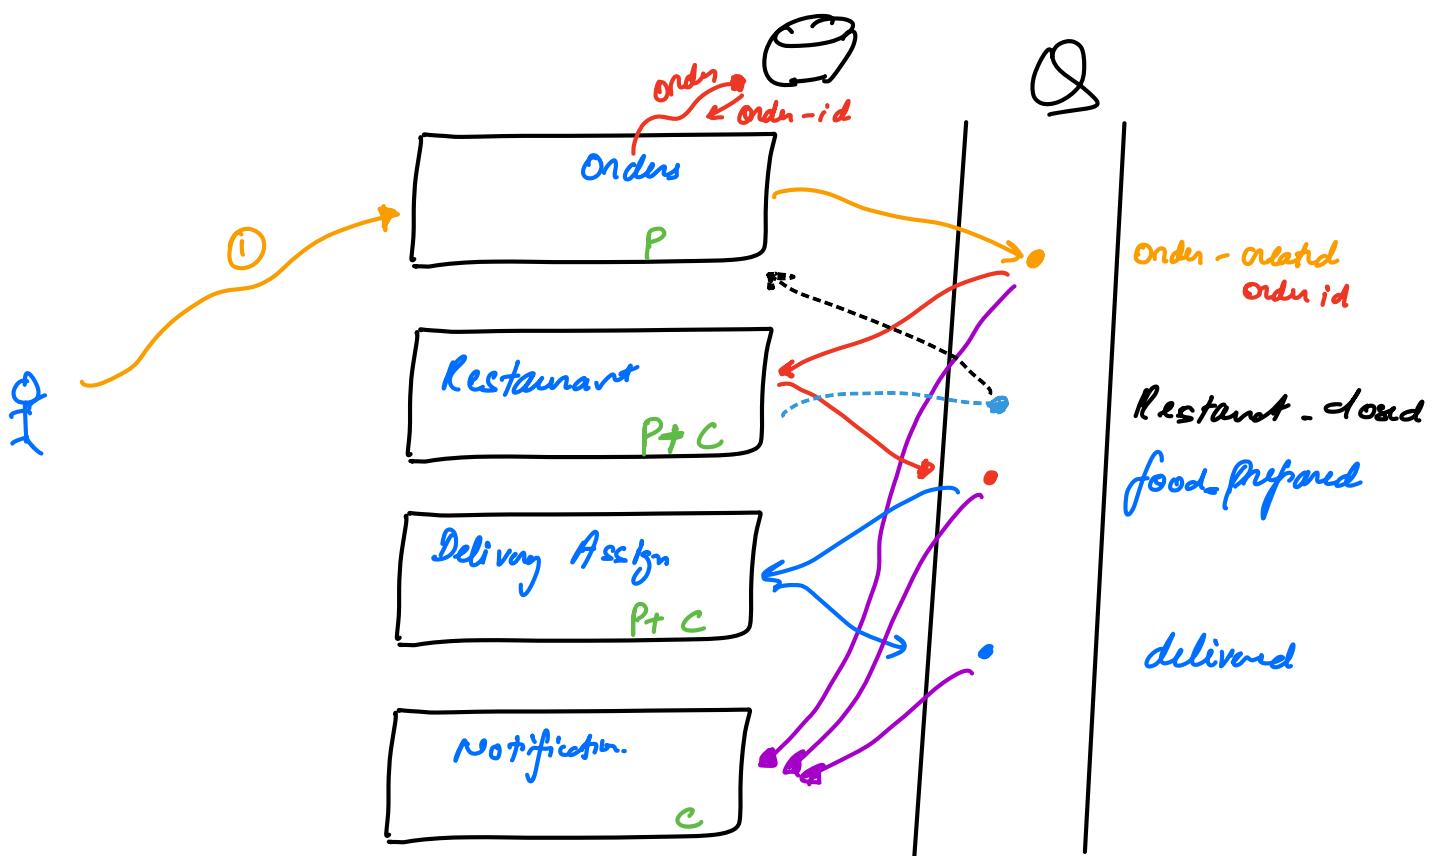
Inherently Distributed — no single orchestrator.

Once a service has done its job it will

Communicate to others that a task has been done.

ANNOUNCE
& forget

— I don't care who listens to it
[my job was to announce]



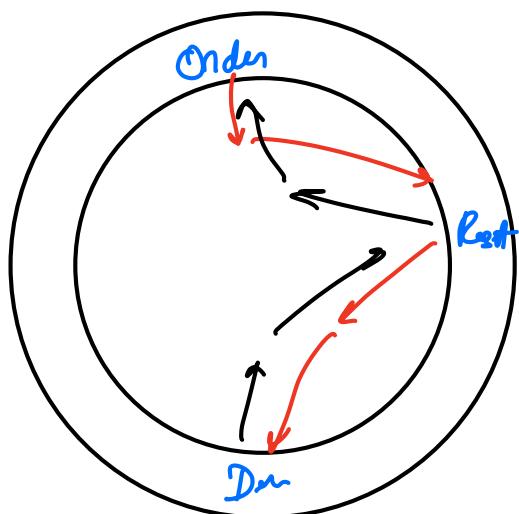
In choreography - each service is (usually) both a producer & a consumer.

- consume tasks that it wants to do
- perform task
- announce task completion.



How to handle failure in choreography

- each service must also announce failure
- each service must consume failures of the immediately next service



Dead Letter Queue

- if a service fails to do a task it might want to retry before announcing failure.
- if all services fail you do NOT want to lose this req completely, you also do NOT want to keep attempting

- drop this rig into a separate topic in Msg Q
 - ↓
 - Dead Lett. Q

Choreography vs Orchestration

choreography is generally better than orchestration.

- ↳ things are more reliable
- ↳ msgs cannot be lost ∵ Kafka
- ↳ if a server fails - no issue ∵ msg can be consumed again.

Cons of choreography

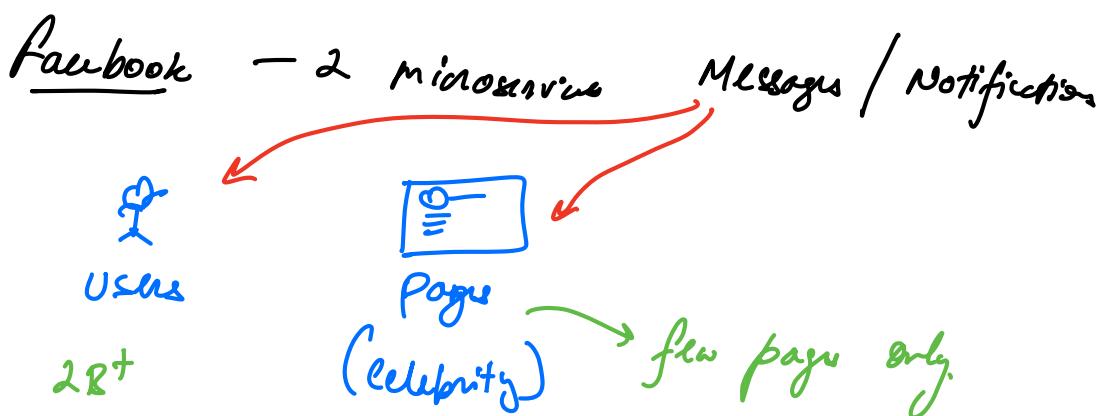
- ① high latency — ∵ inherently event driven
- ② observability is harder — need to have dedicated infrastructure for observability.

10:55 → 11:05

CQRS

(Command / Query Request Segregation)

treat 'commands' (create/update)
del separate from 'queries' (Read)



every time a msg is sent (to user/page) a notification is triggered.

find 1 - How many notifications does a page get on avg? **Notif**
2 - " " " messages " " " " " msg.

3 - what pages get more messages but less notifications.

to ans this (3) you need to join the data b/w
Msg & notification service.

join Small portion of overall data

~~ask~~ ask one service (msg) for this data

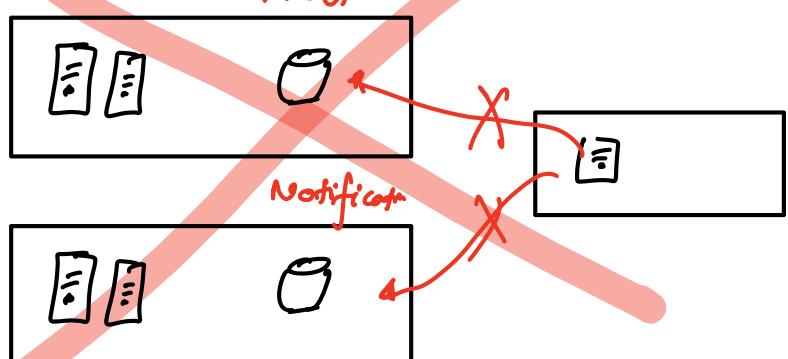
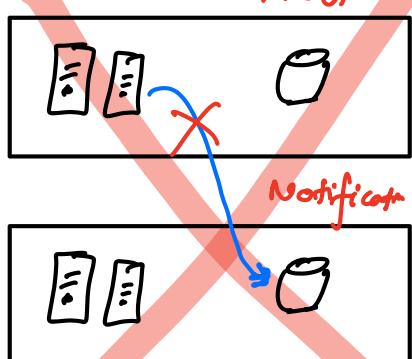
msg service will have its own DB
read the DB of notification service X violation
join

Separation?

return.

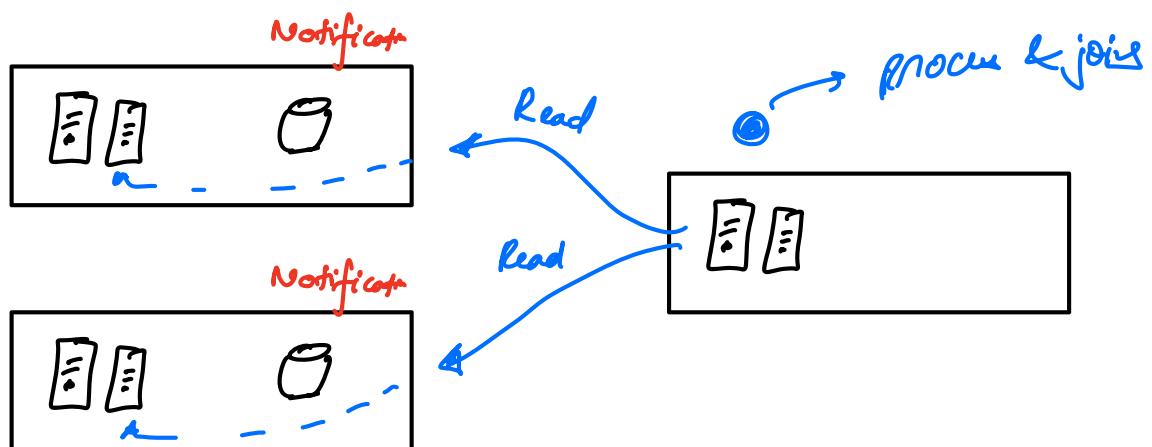
concerns.

~~(2)~~ 3rd service that queries db & both msg & Notify



③

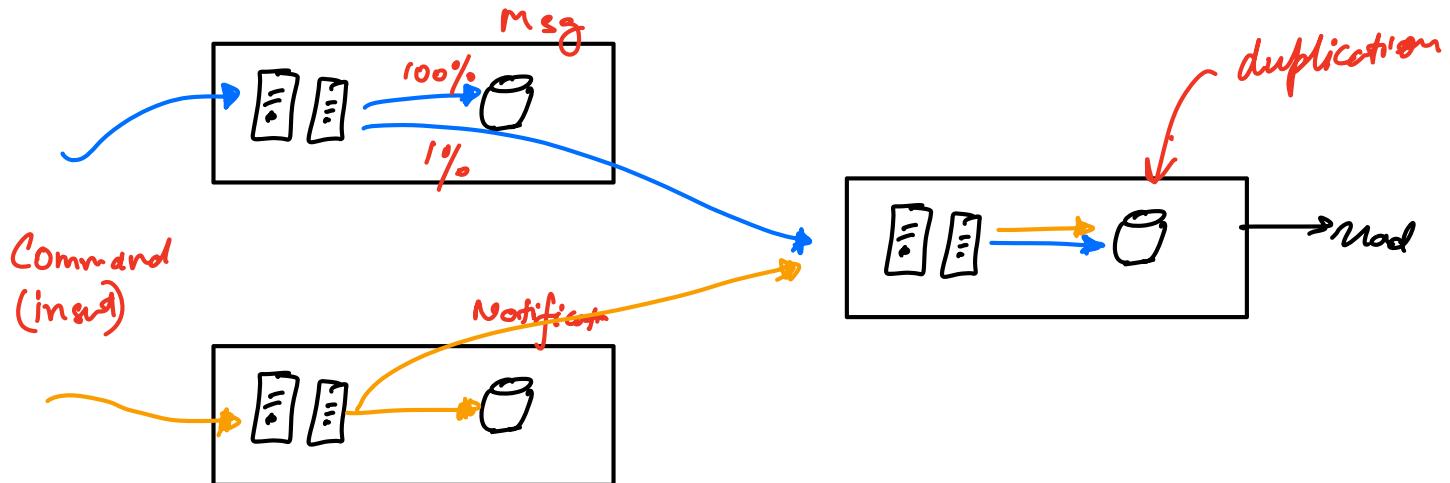
Anti
Pattern



④

CQRS — segregate reads & writes

any reads for fm joint data go to the other service



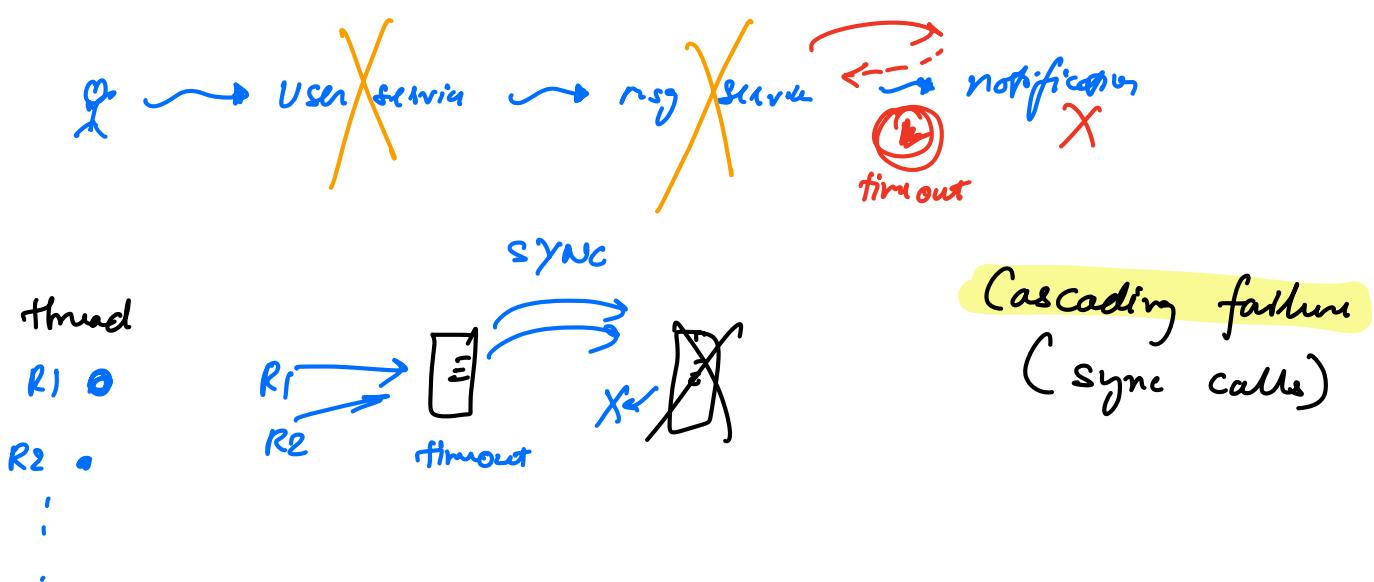
CQRS — instead of fetching data from multiple services A, B, C...

You push data for multiple services into another service X so that any read queries for the joined data of ABC can go to X

Circuit Breaker Pattern

multiple microservices — have dependencies

even more relevant when we are making sync calls



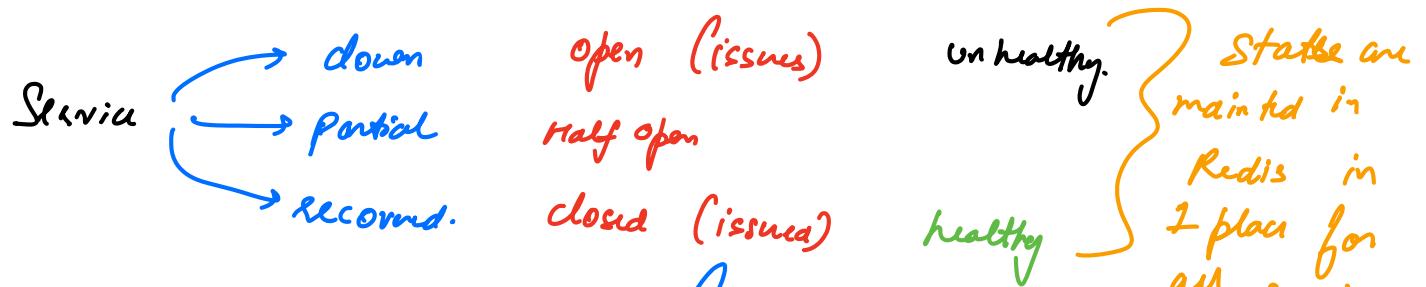
when notification service is "partially fixed"

∴ if the piled up messages it will be slow to process new requests

Thundering Herd

Error

- ↳ Regular - Okay - Service is healthy
error was because of other reasons
auth fail / rate limit / 404
- ↳ Non-functional
↳ bad - Service is unhealthy. - serious issue



- constantly monitoring metrics (# of req successfully processed / min)
 - By default each microservice starts off in closed state
 - if suddenly the metrics fall - open unhealthy.
 - any service dependent on unhealthy service will stop sending additional requests
 - eventually this failed service will be fixed. - Half Open
 - monitor → fail → unhealthy
 - fix → healthy
- calling service will only send a % of its req

